


基于Faster R-CNN进行目标检测



运行

更新于 2020-08-26 09:59:10 3 2

目标检测是计算机视觉领域的重要分支，主要目标是从静态图片，或视频序列中检测并定位感兴趣的目標物体，需要确定目标的类别和位置。本案例中，我们将采用Faster R-CNN模型，使用PyTorch，检测ShipsPascalVOC数据集中，船只在图像中的位置。



目录

- [1. 数据集简介](#)
- [2. 划分数据集](#)
 - [2.1 分析数据集](#)
 - [2.2 划分训练集和测试集](#)
- [3. 建立Faster R-CNN模型](#)
 - [3.1 读取预训练模型](#)
 - [3.2 构建Faster R-CNN的网络结构](#)
- [4. 训练模型](#)
- [5. 评价模型分类效果](#)
- [6. 总结](#)

1 数据集简介

本案例采用的ShipsPascalVOC数据集大小为121MB，包含621张 png 格式的图像，每张图片的大小不完全相同，基本在 400×300 像素左右。

图像中包含不同大小、不同颜色、不同形态和不同关照下的船只。数据集中已经对船只的位置用边框进行了标注，并在 `shiplabels.csv` 文件中注明了每个边框的坐标，便于进行目标识别训练。



2 划分数据集

2.1 分析数据集

在读取数据集之前，我们需要加载案例中使用到的文件和库。首先导入 `utils.py` 、`transforms.py` 等文件，文件中定义了很多需要用到的辅助函数。如 `transform.py` 中定义了水平翻转函数，可以在接下来对训练集进行水平翻转，达到图像增强的效果。

```
# 导入包含辅助函数的python文件
%%bash
git clone https://github.com/pytorch/vision.git
cd vision
git checkout v0.3.0
cp references/detection/utils.py ../

cp references/detection/transforms.py ../

cp references/detection/coco_eval.py ../
cp references/detection/engine.py ../
cp references/detection/coco_utils.py ../
```

```
fatal: destination path 'vision' already exists and is not an empty directory.
HEAD is now at be37608 version check against PyTorch's CUDA version
```

接下来导入需要使用的库。

```
# 导入库
import zipfile
```

```

import numpy as np
import torch
import torch.utils.data
from PIL import Image
import pandas as pd
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from engine import train_one_epoch, evaluate
import utils
import transforms as T
from PIL import Image
import pycocotools
import pandas as pd
import os
import torchvision

```

加载好需要的文件和库后，首先读取ShipsPascalVOC数据集。

```

# 解压zip格式的数据集
zipName = "/content/ShipsPascalVOC.zip"
fileZip = zipfile.ZipFile(zipName)
fileZip.extractall()
fileZip.close()

```

接下来读取 `shiplabels.csv`，其中包含了图像信息。下列表格解读了 `shiplabels.csv` 文件中各列的含义。

列名	含义
filename	图像的名称
width	图像的宽，和 height 一起表示图像的像素大小
height	图像的高
class	图像中包含需要检测的物体的类别，本案例中只检测船只一种物体，故该列值为"boat"
xmin	边框左下角的横坐标，确定边框的位置
ymin	边框左下角的纵坐标
xmax	边框右上角的横坐标
ymax	边框右上角的纵坐标

```

# 改变当前工作目录到指定的路径
os.chdir("/content")
# 读取边框数据
labels = pd.read_csv("/content/shiplabels.csv")
labels.head()

```

	filename	width	height	class	xmin	ymin	xmax	ymax
0	boat0.png	466	418	boat	282	205	333	273

1	boat1.png	400	362	boat	194	128	266	185
2	boat1.png	400	362	boat	169	170	189	187
3	boat10.png	400	243	boat	236	19	249	32
4	boat10.png	400	243	boat	244	68	253	83

接下来我们通过定义函数，在图像上展示 `shiplabels.csv` 中边框的信息。

```
# 返回图像中物体的边框, 每个边框包含四个参数, 分别是左下角和右上角的x轴和y轴的坐标
def parse_one_annot(path_to_data_file, filename):
    data = pd.read_csv(path_to_data_file)
    boxes_array = data[data["filename"] == filename][["xmin", "ymin",
                                                       "xmax", "ymax"]].values

    return boxes_array
```

以"boat10.png"为例，我们读取 `shiplabels.csv` 中的数据，画出图像中包含船只的边框。

```
# 画出"boat10.png"图像
img = cv2.imread("/content/images/boat10.png")
img_rgb = img[:, :, ::-1] # 将图像转换为RGB
plt.imshow(img_rgb)
plt.show()
```



读取 `shiplabels.csv` 中"boat10.png"的边框数据，因为上图中有8个白色船只，都是需要检测到的物体。因此返回 8×4 的数组，表示图中8个船只的边框位置。

```
parse_one_annot("/content/shiplabels.csv", "boat10.png")

array([[236, 19, 249, 32],
       [244, 68, 253, 83],
```

```
[276, 33, 290, 52],
[341, 6, 351, 22],
[340, 58, 350, 73],
[327, 103, 341, 124],
[266, 117, 278, 133],
[384, 115, 394, 131]])
```

如下图，将边框的坐标展现在图像上，可以看到每个需要检测的船只都被红色边框标识出来。

```
image = Image.open("/content/images/boat10.png")
draw = ImageDraw.Draw(image)
coordinates = parse_one_annot("/content/shiplabels.csv", "boat10.png")
for i in coordinates:
    draw.rectangle([(i[0], i[1]), (i[2], i[3])], outline="red")
image
```



我们再定义一个函数，让图像和 `shiplabels.csv` 中的边框数据一一匹配，返回图像中边框的面积、坐标、边框重合度以及边框的个数。

```
# 将图像和csv文件中边框坐标的数据一一对应
class BoatDataset(torch.utils.data.Dataset):
    def __init__(self, root, data_file, transforms=None):
        self.root = root
        self.transforms = transforms

        # 读取图片名
        self.imgs = sorted(os.listdir(os.path.join(root, "images")))
        self.path_to_data_file = data_file

    # 加载图片数据和对应的边框位置
    def __getitem__(self, idx):
        # 加载图片
        img_path = os.path.join(self.root, "images", self.imgs[idx])
        img = Image.open(img_path).convert("RGB")
```

```

# 读取边框位置
box_list = parse_one_annot(self.path_to_data_file, self.imgs[idx])
# 将边框位置转化为Tensor格式
boxes = torch.as_tensor(box_list, dtype=torch.float32)

# 图像中包含的物体的个数
num_objs = len(box_list)
# 需要检测的物体只有一个类别：案例中是船只
labels = torch.ones((num_objs,), dtype=torch.int64)
image_id = torch.tensor([idx])
# 每个边框的面积大小
area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:, 0])
# 假设所有的边框都不是重合度很高的(拥挤的)
iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

# 变量target中包含图像中边框的面积area、边框的坐标boxes、图片的名称image_id、边框重合度(拥挤度)iscrowd、边框的个数labels
target = {}
target["boxes"] = boxes
target["labels"] = labels
target["image_id"] = image_id
target["area"] = area
target["iscrowd"] = iscrowd

# 图像变换
if self.transforms is not None:
    img, target = self.transforms(img, target)
return img, target

# 图像的个数
def __len__(self):
    return len(self.imgs)

```

我们仍以"boat10.png"图像为例，看 `BoatDataset` 的输出结果的形式。

```

# 以上一张图像boat10.png为例，看输出结果
dataset = BoatDataset(root= '/content', data_file= "/content/shiplabels.csv")
dataset.__getitem__(2)

```

```

(<PIL.Image.Image image mode=RGB size=400x243 at 0x7FDE225A2390>,
{'area': tensor([169., 135., 266., 160., 150., 294., 192., 160.]),
'boxes': tensor([[236., 19., 249., 32.],
                  [244., 68., 253., 83.],
                  [276., 33., 290., 52.],
                  [341., 6., 351., 22.],
                  [340., 58., 350., 73.],
                  [327., 103., 341., 124.],
                  [266., 117., 278., 133.],
                  [384., 115., 394., 131.]]),
'image_id': tensor([2]),
'iscrowd': tensor([0, 0, 0, 0, 0, 0, 0, 0]),
'labels': tensor([1, 1, 1, 1, 1, 1, 1, 1])})

```


2.2 划分训练集和测试集

接下来，我们随机抽取数据集中的图像，划分数据集。划分比例为，训练集：测试集 = 7 : 3。故测试集有 $621 \times 0.3 = 186$ 张图像，训练集有 $621 - 186 = 435$ 张图像。

在划分数据集之前，我们需要定义一个函数，对图像随机进行水平翻转的处理，已对训练集进行图像增强。图像增强可以增加模型的鲁棒性，有利于模型识别各种场景的图像。同时需要注意的是，不用对测试集进行图像增强。

```
def get_transform(train):
    transforms = []
    # 将图像数据转换为Tensor格式
    transforms.append(T.ToTensor())
    if train:
        # 在训练过程中，对训练集随机抽取数据，进行图像和边框的水平翻转，进行图像增强
        transforms.append(T.RandomHorizontalFlip(0.5))
    return T.Compose(transforms)
```

```
dataset = BoatDataset(root= "/content",
                      data_file= "/content/shiplabels.csv",
                      transforms = get_transform(train=True))
dataset_test = BoatDataset(root= "/content",
                           data_file= "/content/shiplabels.csv",
                           transforms = get_transform(train=False))

# 分割训练集和测试集
# 设置随机种子
torch.manual_seed(1)

# 随机抽取435张测试集，186张测试集
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:186])
dataset_test = torch.utils.data.Subset(dataset_test, indices[186:])
```

接下来，我们将训练集和测试集转换为 `DataLoader` 格式的数据，方便在模型训练和测试的时候，按照批次(batch)加载数据。

```
# 保存为训练集和测试集
data_loader = torch.utils.data.DataLoader(
    dataset, batch_size=2, shuffle=True, num_workers=4,
    collate_fn=utils.collate_fn)
data_loader_test = torch.utils.data.DataLoader(
    dataset_test, batch_size=1, shuffle=False, num_workers=4,
    collate_fn=utils.collate_fn)
```

```
print("数据集共包含：{} 张图像；训练集：{} 张；测试集：{} 张".format(len(indices), len(dataset), len(dataset_test)))
```

数据集共包含：621 张图像；训练集：435 张；测试集：186 张

3 建立Faster R-CNN模型

接下来我们将构建Faster R-CNN模型。

3.1 读取预训练模型

为了节省模型训练时间，并且因为本案例采用的数据集较小，为了得到更好的训练结果，我们通过迁移学习训练模型。下面我们通过定义函数，加载在COCO数据集上训练好的Faster R-CNN模型。

因为在进行迁移学习的时候，如果需要训练的数据集较小，且与原数据集的相似程度不高，我们需要微调Faster R-CNN模型较前面的层。因为前面的层所提取的特征更具有一般性，后面的层则更加具体，更具有原始数据集的特征。

```
def get_model(num_classes):  
    # 加载在COCO数据集上进行预训练得到的目标检测模型  
    model = torchvision.models.detection.fasterrcnn_resnet50_fpn(pretrained=True)  
  
    # 获取输入分类器的特征的数量  
    in_features = model.roi_heads.box_predictor.cls_score.in_features  
  
    # 修改预训练模型较前面的层  
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)  
    return model
```

3.2 构建Faster R-CNN的网络结构

下面打印出了我们采用的Faster R-CNN的网络结构。

```
# 判断是否用GPU训练模型  
torch.cuda.is_available()  
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')  
  
# 输出 Faster R-CNN 网络的结构  
# 数据集中只包含两类图像；包含船只的图像、不包含船只的图像  
num_classes = 2  
  
# 用辅助函数获得预训练模型
```



```
model = get_model(num_classes)
```

```
# 用合适的设备训练模型
```

```
model.to(device)
```

```
FasterRCNN(
  (transform): GeneralizedRCNNTransform(
    Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    Resize(min_size=(800,), max_size=1333, mode='bilinear')
  )
  (backbone): BackboneWithFPN(
    (body): IntermediateLayerGetter(
      (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3,
3), bias=False)
      (bn1): FrozenBatchNorm2d(64)
      (relu): ReLU(inplace=True)
      (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, cei
l_mode=False)
      (layer1): Sequential(
        (0): Bottleneck(
          (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=Fals
e)
          (bn1): FrozenBatchNorm2d(64)
          (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(
1, 1), bias=False)
          (bn2): FrozenBatchNorm2d(64)
          (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=Fal
se)
          (bn3): FrozenBatchNorm2d(256)
          (relu): ReLU(inplace=True)
          (downsample): Sequential(
            (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=Fals
e)
            (1): FrozenBatchNorm2d(256)
          )
        )
      )
      (1): Bottleneck(
        (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=Fal
se)
        (bn1): FrozenBatchNorm2d(64)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(
1, 1), bias=False)
        (bn2): FrozenBatchNorm2d(64)
        (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=Fal
se)
        (bn3): FrozenBatchNorm2d(256)
        (relu): ReLU(inplace=True)
      )
      (2): Bottleneck(
        (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=Fal
se)
        (bn1): FrozenBatchNorm2d(64)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(
1, 1), bias=False)
        (bn2): FrozenBatchNorm2d(64)
```

```

(conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    se)
        (bn3): FrozenBatchNorm2d(256)
        (relu): ReLU(inplace=True)
    )
)
(layer2): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    lse)
        (bn1): FrozenBatchNorm2d(128)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding
=(1, 1), bias=False)
        (bn2): FrozenBatchNorm2d(128)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    lse)
        (bn3): FrozenBatchNorm2d(512)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
          e)
              (1): FrozenBatchNorm2d(512)
          )
        )
    )
    (1): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      lse)
          (bn1): FrozenBatchNorm2d(128)
          (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1), bias=False)
          (bn2): FrozenBatchNorm2d(128)
          (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      lse)
          (bn3): FrozenBatchNorm2d(512)
          (relu): ReLU(inplace=True)
      )
    )
    (2): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      lse)
          (bn1): FrozenBatchNorm2d(128)
          (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1), bias=False)
          (bn2): FrozenBatchNorm2d(128)
          (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      lse)
          (bn3): FrozenBatchNorm2d(512)
          (relu): ReLU(inplace=True)
      )
    )
    (3): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      lse)
          (bn1): FrozenBatchNorm2d(128)
          (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1), bias=False)
          (bn2): FrozenBatchNorm2d(128)

```

```

(conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    lse)
        (bn3): FrozenBatchNorm2d(512)
        (relu): ReLU(inplace=True)
    )
)
(layer3): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    lse)
        (bn1): FrozenBatchNorm2d(256)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
        (bn2): FrozenBatchNorm2d(256)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    else)
        (bn3): FrozenBatchNorm2d(1024)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
          (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
          se)
              (1): FrozenBatchNorm2d(1024)
          )
        )
    (1): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      else)
          (bn1): FrozenBatchNorm2d(256)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (bn2): FrozenBatchNorm2d(256)
          (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      else)
          (bn3): FrozenBatchNorm2d(1024)
          (relu): ReLU(inplace=True)
      )
    (2): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      else)
          (bn1): FrozenBatchNorm2d(256)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (bn2): FrozenBatchNorm2d(256)
          (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      else)
          (bn3): FrozenBatchNorm2d(1024)
          (relu): ReLU(inplace=True)
      )
    (3): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      else)
          (bn1): FrozenBatchNorm2d(256)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (bn2): FrozenBatchNorm2d(256)

```

```

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=F
else)
    (bn3): FrozenBatchNorm2d(1024)
    (relu): ReLU(inplace=True)
)
(4): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=F
else)
    (bn1): FrozenBatchNorm2d(256)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(256)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=F
else)
    (bn3): FrozenBatchNorm2d(1024)
    (relu): ReLU(inplace=True)
)
(5): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=F
else)
    (bn1): FrozenBatchNorm2d(256)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(256)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=F
else)
    (bn3): FrozenBatchNorm2d(1024)
    (relu): ReLU(inplace=True)
)
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=F
else)
    (bn1): FrozenBatchNorm2d(512)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding
=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(512)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=F
else)
    (bn3): FrozenBatchNorm2d(2048)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=Fa
lse)
      (1): FrozenBatchNorm2d(2048)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=F
else)
    (bn1): FrozenBatchNorm2d(512)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1), bias=False)
    (bn2): FrozenBatchNorm2d(512)

```

```

        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=F
    else)
        (bn3): FrozenBatchNorm2d(2048)
        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=F
    else)
        (bn1): FrozenBatchNorm2d(512)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1), bias=False)
        (bn2): FrozenBatchNorm2d(512)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=F
    else)
        (bn3): FrozenBatchNorm2d(2048)
        (relu): ReLU(inplace=True)
    )
    )
    )
    (fpn): FeaturePyramidNetwork(
        (inner_blocks): ModuleList(
            (0): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
            (1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
            (2): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
            (3): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
        )
        (layer_blocks): ModuleList(
            (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
            (1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
            (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        )
        (extra_blocks): LastLevelMaxPool()
    )
    )
    (rpn): RegionProposalNetwork(
        (anchor_generator): AnchorGenerator()
        (head): RPNHead(
            (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
            (cls_logits): Conv2d(256, 3, kernel_size=(1, 1), stride=(1, 1))
            (bbox_pred): Conv2d(256, 12, kernel_size=(1, 1), stride=(1, 1))
        )
    )
    (roi_heads): RoIHeads(
        (box_roi_pool): MultiScaleRoIAlign()
        (box_head): TwoMLPHead(
            (fc6): Linear(in_features=12544, out_features=1024, bias=True)
            (fc7): Linear(in_features=1024, out_features=1024, bias=True)
        )
        (box_predictor): FastRCNNPredictor(

```

```

        (cls_score): Linear(in_features=1024, out_features=2, bias=True)
        (bbox_pred): Linear(in_features=1024, out_features=8, bias=True)
    )
)
)

```

同时, 我们还需要为模型选择优化算法和学习率调整方式。

这里我们选择随机梯度下降(SGD)算法作为参数优化算法。并将学习率的调整方式设定为: 一次epoch代表对所有样本进行一次训练, 每进行3次epoch, 学习率下降10倍。

```

# 选择SGD优化器算法
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005,
                             momentum=0.9, weight_decay=0.0005)

# 确定调整学习率的方式: 每进行3次epoch降低10倍学习率
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                                  step_size=3,
                                                  gamma=0.1)

```

4 训练模型

在定义好上述函数后, 我们调用这些函数, 在ShipsPascalVOC数据集上进行模型的训练。

```

# 训练5次epoch
num_epochs = 5
for epoch in range(num_epochs):
    # 打印每次训练的结果
    train_one_epoch(model, optimizer, data_loader, device, epoch,
                    print_freq=100)
    # 更新学习率
    lr_scheduler.step()
    # 看模型在测试集上训练的效果
    evaluate(model, data_loader_test, device=device)

```

```

Epoch: [0] [ 0/218] eta: 0:05:34 lr: 0.000028 loss: 0.0443 (0.0443) loss
_classifier: 0.0136 (0.0136) loss_box_reg: 0.0301 (0.0301) loss_objectness:
0.0001 (0.0001) loss_rpn_box_reg: 0.0006 (0.0006) time: 1.5324 data: 0.2911
max mem: 4598
Epoch: [0] [100/218] eta: 0:02:38 lr: 0.002330 loss: 0.0461 (0.1145) loss
_classifier: 0.0149 (0.0295) loss_box_reg: 0.0287 (0.0625) loss_objectness:
0.0002 (0.0043) loss_rpn_box_reg: 0.0007 (0.0182) time: 1.2581 data: 0.0109
max mem: 4598
Epoch: [0] [200/218] eta: 0:00:23 lr: 0.004632 loss: 0.0519 (0.1061) loss
_classifier: 0.0159 (0.0271) loss_box_reg: 0.0313 (0.0587) loss_objectness:
0.0001 (0.0038) loss_rpn_box_reg: 0.0013 (0.0165) time: 1.3309 data: 0.0094
max mem: 4598
Epoch: [0] [217/218] eta: 0:00:01 lr: 0.005000 loss: 0.0607 (0.1039) loss

```



```

Epoch: [0] [ 0/186] eta: 0:01:59 model_time: 0.3562 (0.3562) evaluator_time:
_classifier: 0.0189 (0.0267) loss_box_reg: 0.0406 (0.0583) loss_objectness:
0.0001 (0.0036) loss_rpn_box_reg: 0.0007 (0.0153) time: 1.2824 data: 0.0090
max mem: 4598
Epoch: [0] Total time: 0:04:49 (1.3264 s / it)
creating index...
index created!
Test: [ 0/186] eta: 0:01:59 model_time: 0.3562 (0.3562) evaluator_time:
0.0027 (0.0027) time: 0.6401 data: 0.2743 max mem: 4598
Test: [100/186] eta: 0:00:27 model_time: 0.2796 (0.3011) evaluator_time:
0.0013 (0.0049) time: 0.2983 data: 0.0050 max mem: 4598
Test: [185/186] eta: 0:00:00 model_time: 0.2841 (0.3027) evaluator_time:
0.0013 (0.0050) time: 0.3152 data: 0.0048 max mem: 4598
Test: Total time: 0:00:58 (0.3157 s / it)
Averaged stats: model_time: 0.2841 (0.3027) evaluator_time: 0.0013 (0.0050)
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.46
5
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.83
7
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.43
2
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.28
2
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.70
2
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.85
9
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.25
0
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.39
0
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.51
0
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.35
7
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.75
3
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.88
5
Epoch: [1] [ 0/218] eta: 0:06:04 lr: 0.005000 loss: 0.0282 (0.0282) loss
_classifier: 0.0145 (0.0145) loss_box_reg: 0.0136 (0.0136) loss_objectness:
0.0000 (0.0000) loss_rpn_box_reg: 0.0002 (0.0002) time: 1.6724 data: 0.3879
max mem: 4598
Epoch: [1] [100/218] eta: 0:02:37 lr: 0.005000 loss: 0.0528 (0.1240) loss
_classifier: 0.0150 (0.0322) loss_box_reg: 0.0360 (0.0668) loss_objectness:
0.0001 (0.0040) loss_rpn_box_reg: 0.0009 (0.0210) time: 1.3174 data: 0.0096
max mem: 4598
Epoch: [1] [200/218] eta: 0:00:23 lr: 0.005000 loss: 0.0608 (0.1183) loss
_classifier: 0.0176 (0.0298) loss_box_reg: 0.0404 (0.0641) loss_objectness:
0.0002 (0.0041) loss_rpn_box_reg: 0.0008 (0.0203) time: 1.2939 data: 0.0097
max mem: 4598
Epoch: [1] [217/218] eta: 0:00:01 lr: 0.005000 loss: 0.0528 (0.1152) loss
_classifier: 0.0143 (0.0294) loss_box_reg: 0.0313 (0.0626) loss_objectness:

```

```

0.0001 (0.0039) loss_rpn_box_reg: 0.0011 (0.0193) time: 1.2486 data: 0.0092
max mem: 4598
Epoch: [1] Total time: 0:04:47 (1.3167 s / it)
creating index...
index created!
Test: [ 0/186] eta: 0:02:02 model_time: 0.3643 (0.3643) evaluator_time:
0.0025 (0.0025) time: 0.6571 data: 0.2886 max mem: 4598
Test: [100/186] eta: 0:00:27 model_time: 0.2838 (0.3016) evaluator_time:
0.0013 (0.0054) time: 0.3000 data: 0.0050 max mem: 4598
Test: [185/186] eta: 0:00:00 model_time: 0.2843 (0.3032) evaluator_time:
0.0015 (0.0053) time: 0.3152 data: 0.0050 max mem: 4598
Test: Total time: 0:00:58 (0.3165 s / it)
Averaged stats: model_time: 0.2843 (0.3032) evaluator_time: 0.0015 (0.0053)
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.49
1
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.85
7
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.49
9
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.31
8
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.71
9
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.82
2
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.25
4
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.40
6
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.54
5
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.40
9
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.77
0
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.85
6
Epoch: [2] [ 0/218] eta: 0:06:37 lr: 0.000500 loss: 0.2067 (0.2067) loss
_classifier: 0.0668 (0.0668) loss_box_reg: 0.1217 (0.1217) loss_objectness:
0.0147 (0.0147) loss_rpn_box_reg: 0.0035 (0.0035) time: 1.8238 data: 0.4054
max mem: 4598
Epoch: [2] [200/218] eta: 0:00:23 lr: 0.000500 loss: 0.0442 (0.0993) loss
_classifier: 0.0152 (0.0261) loss_box_reg: 0.0304 (0.0543) loss_objectness:
0.0003 (0.0028) loss_rpn_box_reg: 0.0004 (0.0161) time: 1.3852 data: 0.0094
max mem: 4598
Epoch: [2] [217/218] eta: 0:00:01 lr: 0.000500 loss: 0.0395 (0.0962) loss
_classifier: 0.0123 (0.0254) loss_box_reg: 0.0236 (0.0530) loss_objectness:
0.0001 (0.0026) loss_rpn_box_reg: 0.0003 (0.0152) time: 1.2655 data: 0.0091
max mem: 4598
Epoch: [2] Total time: 0:04:48 (1.3251 s / it)
creating index...
index created!

```

```

Test: [ 0/186] eta: 0:01:52 model_time: 0.3477 (0.3477) evaluator_time:
0.0026 (0.0026) time: 0.6025 data: 0.2507 max mem: 4598
Test: [100/186] eta: 0:00:27 model_time: 0.2822 (0.3025) evaluator_time:
0.0013 (0.0051) time: 0.3010 data: 0.0050 max mem: 4598
Test: [185/186] eta: 0:00:00 model_time: 0.2873 (0.3045) evaluator_time:
0.0014 (0.0052) time: 0.3192 data: 0.0049 max mem: 4598
Test: Total time: 0:00:59 (0.3179 s / it)
Averaged stats: model_time: 0.2873 (0.3045) evaluator_time: 0.0014 (0.0052)
Accumulating evaluation results...
DONE (t=0.04s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.50
3
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.85
6
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.50
2
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.32
2
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.74
1
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.86
3
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.26
4
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.41
7
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.55
7
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.41
2
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.79
3
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.89
4
Epoch: [3] [ 0/218] eta: 0:06:02 lr: 0.000500 loss: 0.0459 (0.0459) loss
_classifier: 0.0175 (0.0175) loss_box_reg: 0.0274 (0.0274) loss_objectness:
0.0001 (0.0001) loss_rpn_box_reg: 0.0009 (0.0009) time: 1.6618 data: 0.4468
max mem: 4598
Epoch: [3] [100/218] eta: 0:02:38 lr: 0.000500 loss: 0.0347 (0.0886) loss
_classifier: 0.0125 (0.0235) loss_box_reg: 0.0225 (0.0466) loss_objectness:
0.0001 (0.0026) loss_rpn_box_reg: 0.0003 (0.0159) time: 1.2835 data: 0.0103
max mem: 4598
Epoch: [3] [200/218] eta: 0:00:24 lr: 0.000500 loss: 0.0302 (0.0910) loss
_classifier: 0.0107 (0.0242) loss_box_reg: 0.0190 (0.0505) loss_objectness:
0.0001 (0.0027) loss_rpn_box_reg: 0.0003 (0.0135) time: 1.3463 data: 0.0099
max mem: 4598
Epoch: [3] [217/218] eta: 0:00:01 lr: 0.000500 loss: 0.0279 (0.0902) loss
_classifier: 0.0103 (0.0240) loss_box_reg: 0.0161 (0.0498) loss_objectness:
0.0001 (0.0026) loss_rpn_box_reg: 0.0003 (0.0137) time: 1.2636 data: 0.0104
max mem: 4598
Epoch: [3] Total time: 0:04:50 (1.3336 s / it)
creating index...
index created!
Test: [ 0/186] eta: 0:01:52 model time: 0.3774 (0.3774) evaluator time:

```

```

0.0026 (0.0026) time: 0.6055 data: 0.2238 max mem: 4598
Test: [100/186] eta: 0:00:27 model_time: 0.2814 (0.3027) evaluator_time:
0.0013 (0.0051) time: 0.3001 data: 0.0049 max mem: 4598
Test: [185/186] eta: 0:00:00 model_time: 0.2854 (0.3044) evaluator_time:
0.0014 (0.0051) time: 0.3177 data: 0.0049 max mem: 4598
Test: Total time: 0:00:59 (0.3172 s / it)
Averaged stats: model_time: 0.2854 (0.3044) evaluator_time: 0.0014 (0.0051)
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.50
8
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.86
6
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.51
2
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.32
7
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.74
0
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.86
9
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.26
2
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.41
7
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.56
1
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.41
9
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.79
4
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.89
3
Epoch: [4] [ 0/218] eta: 0:05:18 lr: 0.000500 loss: 0.0358 (0.0358) loss
_classifier: 0.0121 (0.0121) loss_box_reg: 0.0237 (0.0237) loss_objectness:
0.0000 (0.0000) loss_rpn_box_reg: 0.0001 (0.0001) time: 1.4631 data: 0.2642
max mem: 4598
Epoch: [4] [100/218] eta: 0:02:34 lr: 0.000500 loss: 0.0380 (0.0790) loss
_classifier: 0.0127 (0.0222) loss_box_reg: 0.0249 (0.0452) loss_objectness:
0.0002 (0.0017) loss_rpn_box_reg: 0.0004 (0.0099) time: 1.2797 data: 0.0098
max mem: 4598
Epoch: [4] [200/218] eta: 0:00:23 lr: 0.000500 loss: 0.0368 (0.0926) loss
_classifier: 0.0121 (0.0243) loss_box_reg: 0.0241 (0.0510) loss_objectness:
0.0001 (0.0027) loss_rpn_box_reg: 0.0006 (0.0146) time: 1.3716 data: 0.0104
max mem: 4598
Epoch: [4] [217/218] eta: 0:00:01 lr: 0.000500 loss: 0.0262 (0.0887) loss
_classifier: 0.0103 (0.0235) loss_box_reg: 0.0179 (0.0490) loss_objectness:
0.0001 (0.0025) loss_rpn_box_reg: 0.0003 (0.0136) time: 1.3427 data: 0.0095
max mem: 4598
Epoch: [4] Total time: 0:04:48 (1.3248 s / it)
creating index...
index created!
Test: [ 0/186] eta: 0:01:54 model_time: 0.3720 (0.3720) evaluator_time:
0.0024 (0.0024) time: 0.6182 data: 0.2419 max mem: 4598

```

```

Test: [100/186] eta: 0:00:27 model_time: 0.2812 (0.3038) evaluator_time:
0.0013 (0.0051) time: 0.3015 data: 0.0050 max mem: 4598
Test: [185/186] eta: 0:00:00 model_time: 0.2874 (0.3054) evaluator_time:
0.0014 (0.0052) time: 0.3189 data: 0.0050 max mem: 4598
Test: Total time: 0:00:59 (0.3186 s / it)
Averaged stats: model_time: 0.2874 (0.3054) evaluator_time: 0.0014 (0.0052)
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.50
5
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.85
9
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.51
5
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.32
5
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.73
5
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.86
7
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.26
3
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.41
4
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.55
7
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.41
4
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.78
8
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.89
3

```

因为我们设定训练每一个epoch的过程中，每50张打印一次训练结果。训练结果包括：

在训练集中：

- eta : 每训练50张图像所需要的时间；
- lr : 学习率；
- loss : 损失函数；
- loss_classifier : 分类损失；
- loss_box_reg : 预测边框和真实边框的误差；
- loss_objectness : 边框包含物体的置信度损失；
- loss_rpn_box_reg : 计算RPN的边界损失。

在测试集中：

- Average Precision : 平均准确率；
- IOU : 预测边框和目标区域的交集除以并集；
- IoU=0.50 : 重合比例大于0.5的算正例；

- IoU=0.50:0.95 : 从0.50到0.95每隔0.5取一次作为阈值，计算一次，然后取均值。
- area : 区域大小， small 为小于 32×32 的区域； medium 为 32×32 和 96×96 之间的区域； large 为大于 96×96 的区域；
- maxDets : 最多取目标区域的个数；
- Average Recall : 平均召回率。

因此，从输出结果我们可以看到，模型对较大的物体，即 large 区域的检测准确率较高，在 IoU=0.50:0.95 的情况下达到0.867。同时，在 IoU=0.50 的情况下，模型检测各种尺度的物体，达到0.859的平均准确率。

5 评价模型分类效果

接下来我们观察模型的目标检测效果如何。首先读取模型训练后得到的参数。

```
loaded_model = get_model(num_classes = 2)
loaded_model.load_state_dict(model.state_dict())
```

<All keys matched successfully>

下图展示了模型在随机抽取的一张图像上，对船只进行目标检测的结果。

```
idx = 2
img, _ = dataset_test[idx]
label_boxes = np.array(dataset_test[idx][1]["boxes"])
# 调用训练好的模型的参数，观测其检测物体的效果
loaded_model.eval()
with torch.no_grad():
    prediction = loaded_model([img])
image = Image.fromarray(img.mul(255).permute(1, 2, 0).byte().numpy())
draw = ImageDraw.Draw(image)

# 画出真实边框和预测边框
for elem in range(len(label_boxes)):
    draw.rectangle([(label_boxes[elem][0], label_boxes[elem][1]),
                    (label_boxes[elem][2], label_boxes[elem][3])],
                  outline="green", width=3)
for element in range(len(prediction[0]["boxes"])):
    boxes = prediction[0]["boxes"][element].cpu().numpy()
    score = np.round(prediction[0]["scores"][element].cpu().numpy(),
                      decimals=4)
    if score > 0.8:
        draw.rectangle([(boxes[0], boxes[1]), (boxes[2], boxes[3])],
                      outline="red", width=3)
        draw.text((boxes[0], boxes[1]), text = str(score))
image
```





其中绿色的边框是真实的包含船只的边框，即 `shiplabels.csv` 中该图像对应的数据。红色的边框是由模型预测得到的。红色边框左上角的数字 0.9985 表示真实边框和预测边框的交并比。由于二者的交并比已经达到 0.9985，说明模型在这一张图像上的预测结果很好。

6 总结

本案例介绍了如何使用Faster R-CNN模型对ShipsPascalVOC数据集进行目标检测。虽然数据集图像较少，但通过迁移学习和微调的方法，我们仍获得了较好的精确率。并且通过随机抽取图像进行目标检测，看到模型的预测边框与真实边框重合度较高，检测结果较好。

发表您的讨论

© 2018 CookData 京ICP备1705652	竞赛	关于	产品	服务	帮助	联系
3号-1 (http://www.beian.miit.gov.cn/)	平台	我们 (/	介绍 (/	条款 (/	中心 (/	我们 (/
	(http://	foote	foote	foote	foote	foote
	cookd	r?sub	r?sub	r?sub	r?sub	r?sub
	ata.c	=0)	=0)	=1)	=2)	=3)
	n/com					
	petitio					
	n/)					